

# Address Sanitizer(ASan)

---

## Introduction

---

AddressSanitizer 是一个快速的内存错误探测器。它由一个编译检测模块和一个运行时库组成。它可以检测到一下类型的bug:

- Out-of-bounds accesses to heap, stack and globals
- Use-after-free
- Use-after-return (to some extent)//未测试
- Double-free, invalid free
- Memory leaks (experimental)//未测试

Address Sanitizer重写了malloc和free,是一种非常高效的内存检查方式,使用Address Sanitizer后程序执行仅仅慢了两倍。

AddressSanitizer是一对工具: 首先是编译扩展, 然后是运行库。ASan的运行库会分配一个影子内存: 一个超大块的内存, 用于为每8字节内存记录一字节信息。默认情况下, 所有内存的影子字节会被设置为0, 表明它不能被访问。当内存被分配后, 影子字节设置为其它值(用于记录字的哪些字节被分配、谁分配了它们), 它还会重载分配器来跟踪内存的分配和释放。

这样, 每次发生内存访问时, 运行库将检查相应影子字节的值, 如果访问是不被允许的, ASan就会中止程序的执行: ASan在第一个错误发生时崩溃程序, 它强制程序必须是ASan检查通过的。

编译扩展的主要作用是对内存的每一次访问封装在一个小分支中, 通过检查影子内存的内容确认访问是否允许。不过因为是在编译器中进行处理, 所以它可以访问大量的信息, 比如正在访问什么内存, 变量或结构体成员的布局是怎样的, ...并且它还能可以改变这些。这正是ASan让人眼前一亮的地方: 它可以在全局的变量间或栈上的变量间添加红色区域, 使得对这些变量的错误访问更容易被检查到。

## 用法

使用-fsanitize=address 简单的编译链接你的程序。AdressSanitizer运行时库应当最后被被链接,所以确保在链接的最后一步使用clang(not ld)。当链接共享的库文件的时候, Address Sanitizer 运行时库不会被链接, 因此,-wl,-z,defs可能会引起链接错误(不要和Address Sanitizer一起使用这些命令)。为了获取合理的性能,可以添加-o1或者更高的编译参数。

```
% clang -fsanitize=address main.c
```

**fno-omit-frame-pointer**:可以获得更加详细的信息

If a bug is detected, the program will print an error message to stderr and exit with a non-zero exit code. AddressSanitizer exits on the first detected error. This is by design:

- This approach allows AddressSanitizer to produce faster and smaller generated code (both by ~5%).
- Fixing bugs becomes unavoidable. AddressSanitizer does not produce false alarms. Once a memory corruption occurs, the program is in an inconsistent state, which could lead to confusing results and potentially misleading subsequent reports.

如果你的进程是在OSX 10.10及以后的沙盒环境中运行,你需要设置DYLDINSERTLIBRARIES环境变量,并将它指向和编译器打包在一起的ASan库。

## 符号化输出

为了让AddressSanitizer将其输出符号化为你需要的内容, 你需要将 `ASAN_SYMBOLIZER_PATH` 环境变量设置为为llvm-symbolizer库(或者将llvm-symbolizer加到你的\$PATH中)。

```
% ASAN_SYMBOLIZER_PATH=/usr/local/bin/llvm-symbolizer ./a.out
```

如果这个没有效果,你可以使用一个单独的脚本来离线的解析你的输出结果。(在线解析可以使用 `ASAN_OPTIONS=symbolize=0` 来强制禁用)。

```
% ASAN_OPTIONS=symbolize=0 ./a.out 2> log
% projects/compiler-rt/lib/asan/scripts/asan_symbolize.py / < log | c++filt
```

Note:在OSX中你可能需要运行dsymutil。(Note that on OS X you may need to run dsymutil on your binary to have the file:line info in the AddressSanitizer reports.)

Example:

```
void MyFunction(int nSize);
int main(int argc, char * argv[]) {
    int *a = (int*)malloc(3 * sizeof(int));
    a [4] = 1; //Out-of-bounds accesses to heap, stack and globals
    free(a);
    a [6] = 1; //Use-after-free
    free(a); //Double-free, invalid free
}
```

```

jianghaideMac-mini:Desktop jianghai$ clang -fsanitize=address main.c
jianghaideMac-mini:Desktop jianghai$ ./a.out
=====
==2616==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x6020000eebc at pc 0x000108690ef2 bp 0x7fff5756fbd0 sp 0x7fff5756fbc8
WRITE of size 4 at 0x6020000eebc thread T0
    #0 0x108690ef1 in atos[2617]: [fatal] 'pid_for_task' failed: (os/kern) failure (5) (+0x100000ef1)
==2616==WARNING: Can't write to symbolizer at fd 3
    #1 0x7fff8bbbb5ac in atos[2620]: [fatal] 'pid_for_task' failed: (os/kern) failure (5) (+0x35ac)
    #2 0x0 (<unknown module>)

0x6020000eebc is located 0 bytes to the right of 12-byte region [0x6020000eeb0,0x6020000eebc)
allocated by thread T0 here:
==2616==WARNING: Can't write to symbolizer at fd 3
    #0 0x1086e19c0 in atos[2623]: [fatal] 'pid_for_task' failed: (os/kern) failure (5) (+0x489c0)
==2616==WARNING: Can't write to symbolizer at fd 3
    #1 0x108690e9d in atos[2626]: [fatal] 'pid_for_task' failed: (os/kern) failure (5) (+0x100000e9d)
==2616==WARNING: Can't read from symbolizer at fd 3
==2616==WARNING: Failed to use and restart external symbolizer!
    #2 0x7fff8bbbb5ac in start (/usr/lib/system/libdyld.dylib+0x35ac)
    #3 0x0 (<unknown module>)

SUMMARY: AddressSanitizer: heap-buffer-overflow (/Users/jianghai/Desktop/./a.out+0x100000ef1) in main
Shadow bytes around the buggy address:
 0x1c0400001d80: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001d90: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001da0: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001db0: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001dc0: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
=>0x1c0400001dd0: fa fa fa fa fa fa 00[04]fa fa 00 06 fa fa 00 00
 0x1c0400001de0: fa fa 00 04 fa fa 00 00 fa fa 00 00 fa fa 00 fa
 0x1c0400001df0: fa fa fd fd fa fa fd fd fa fa fd fd fa fa fd fa
 0x1c0400001e00: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001e10: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x1c0400001e20: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
Shadow byte legend (one shadow byte represents 8 application bytes):
Addressable: 00
Partially addressable: 01 02 03 04 05 06 07
Heap left redzone: fa
Heap right redzone: fb
Freed heap region: fd
Stack left redzone: f1
Stack mid redzone: f2
Stack right redzone: f3
Stack partial redzone: f4
Stack after return: f5
Stack use after scope: f8
Global redzone: f9
Global init order: f6
Poisoned by user: f7
Container overflow: fc
Array cookie: ac
Intra object redzone: bb
ASan internal: fe
Left alloca redzone: ca
Right alloca redzone: cb

```